

Compboost

Modular framework for component-wise boosting

Daniel Schalk

June 27, 2018

LMU Munich

Working Group Computational Statistics



Table of Contents

Why Component-Wise Boosting

About Compboost


Small Usecase

Small Benchmark

Next Steps

Why Component-Wise Boosting

Why Component-Wise Boosting?



Why Component-Wise Boosting?

- Inherent (unbiased) feature selection [HHKS11].
- The resulting model is sparse since the important models are selected first.
- The parameters are updated iteratively. Therefore, the parameters are estimated on the fly and can be interpreted due to linearity of the base-learners.
- Efficient model for learning in high-dimensional feature spaces ($p \gg n$).

Most popular package for model-based boosting is `mboost`:

- Huge functionality: Lots of available base-learner and losses.
- Suited to boost more complex learner to do, for instance, survival analyses.
- Extendible with own base-learner and losses.

So, why another boosting implementation?

- Major parts of the algorithm are implemented in R which is not the best choice for expensive algorithms.
- The implementation is hard to debug due to nested scopes and no standard R system like S3 or S4.

About Comboost

About Compboost

- With `mboost` as standard, we want to keep the modular principle of defining custom base-learner and losses.
- Completely written in C++ and exposed by Rcpp to obtain high performance and full memory control.
- R API is written with R6 to provide convenient wrapper.
- Major parts of the `compboost` functionality are unit tested against `mboost` to ensure correct use.

Comboost: Functionality

Main components:

- Base-learner and loss classes.
- Logger class for early stopping and logging mechanisms.

Possible extensions:

- Custom R or C++ base-learner.
- Custom R or C++ loss objects.
- Custom logging and stopping rules via custom losses.

Custom classes can be defined without recompiling the whole package, even when using C++ functions.

Small Usecase

Initializing Model

```
library(compboost)

data(PimaIndiansDiabetes, package = "mlbench")

cboost = Compboost$new(PimaIndiansDiabetes, "diabetes",
  loss = BinomialLoss$new())

cboost$addBaselearner("age", "spline", PSplineBlearner,
  degree = 3, knots = 10, penalty = 2, differences = 2)
cboost$addBaselearner("mass", "linear", PolynomialBlearner,
  degree = 1, intercept = TRUE)
```

Initializing Model

```
cboost$train(2000, trace = FALSE)
cboost

## Component-Wise Gradient Boosting
##
## Trained on PimaIndiansDiabetes with target diabetes
## Number of base-learners: 2
## Learning rate: 0.05
## Iterations: 2000
## Positive class: neg
## Offset: 0.3118
##
## BinomialLoss Loss:
##
## Loss function:  $L(y,x) = \log(1 + \exp(-2yf(x)))$ 
##
##
```

Access Results

```
table(cboost$selected())
```

```
##
```

```
##  age_spline mass_linear
```

```
##      1511      489
```

```
head(cboost$risk())
```

```
## [1] 0.6582 0.6565 0.6549 0.6533 0.6518 0.6503
```

```
str(cboost$coef())
```

```
## List of 3
```

```
## $ age_spline : num [1:14, 1] 4.407 0.913 1.124 -1.044 0.126 ...
```

```
## $ mass_linear: num [1:2, 1] 3.0352 -0.0886
```


```
## $ offset      : num 0.312
```

Plot Results

```
cboost$plot("age_spline", iters = c(100, 500, 1000, 2000))
```

Effect of age_spline

Additive contribution of predictor



Custom Poisson Loss: Definition

```
lossPoisson = function (truth, response) {  
  return (-log(exp(response)^truth * exp(-exp(response)) / gamma(truth + 1)))  
}  
gradPoisson = function (truth, response) {  
  return (exp(response) - truth)  
}  
constInitPoisson = function (truth) {  
  return (log(mean.default(truth)))  
}  
# Define custom loss:  
my.poisson.loss = CustomLoss$new(lossPoisson, gradPoisson, constInitPoisson)
```

Custom Poisson Loss: Train Model

```
data(VonBort, package = "vcd")

cboost = Comboost$new(VonBort, "deaths", loss = my.poisson.loss)
cboost$addBaselearner("year", "spline", PSplineBlearner,
  degree = 3, knots = 10, penalty = 2, differences = 2)
cboost$train(100, trace = FALSE)


mod = mboost(deaths ~ bbs(year, differences = 2, lambda = 2,
  degree = 3, knots = 10), data = VonBort, family = Poisson(),
  control = boost_control(mstop = 100, nu = 0.05))

head(cbind(
  compboost = as.numeric(cboost$coef()[[1]]),
  mboost = as.numeric(coef(mod)[[1]])))

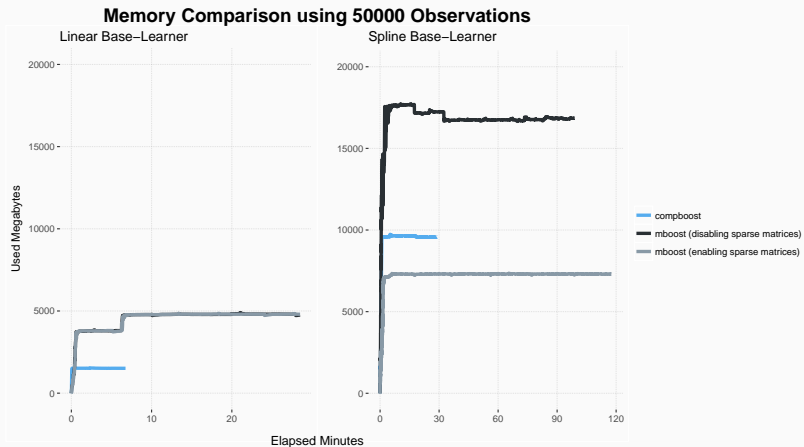
##      compboost  mboost
## [1,]   -1.4093 -1.4093
## [2,]   -0.8975 -0.8975
## [3,]   -0.4040 -0.4040
## [4,]    0.0947  0.0947
## [5,]    0.3384  0.3384
## [6,]    0.1531  0.1531
```


Small Benchmark

Runtime Comparison Against Mboost



Memory Comparison Against Mboost



Next Steps

Next Steps

- Implementing more base-learner and loss functions.
- Support for multiclass classification.
- Parallel computations.
- Using sparse matrices to store, e.g., spline data matrices.

Where to Find

- Developed on *GitHub*:

www.github.com/schalkdaniel/compboost

- Additional resources on the project page:

www.compboost.org

- Bug reports via the issue tracker.
- Contributions are welcome!

Questions?

References



Dirk Eddelbuettel.

Seamless R and C++ Integration with Rcpp.

Springer, New York, 2013.

ISBN 978-1-4614-6867-7.



Dirk Eddelbuettel and Romain François.

Exposing C++ functions and classes with rcpp modules.

Vignette included in R package Rcpp, URL <http://CRAN.R-project.org/package=Rcpp>, 2017.



Torsten Hothorn, Peter Buehlmann, Thomas Kneib, Matthias Schmid, and Benjamin Hofner.

mboost: Model-Based Boosting, 2017.

R package version 2.8-1.



Benjamin Hofner, Torsten Hothorn, Thomas Kneib, and Matthias Schmid.

A framework for unbiased model selection based on boosting.

Journal of Computational and Graphical Statistics, 20(4):956–971, 2011.