

# Compboost

Modular framework for component-wise boosting

---

Daniel Schalk

June 23, 2018

LMU Munich

Working Group Computational Statistics



# Table of contents

1. What is Component-Wise Boosting
2. Compboost vs Other Implementations
3. About Compboost
4. Small Usecase
5. Next Steps

# What is Component-Wise Boosting

---

# Component-Wise Boosting: Terminology

- Loss Function:

$$L : \mathcal{Y} \times \mathcal{X} \rightarrow \mathbb{R}$$

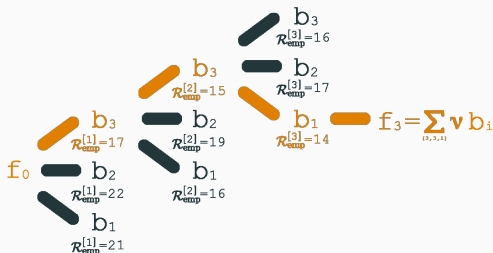
- Empirical Risk:

$$\mathcal{R}_{\text{emp}}(\theta) = \frac{1}{n} \sum_{i=1}^n L\left(y^{(i)}, f(x^{(i)})\right)$$

- Estimated model/parameter at iteration  $m$ :

$$\hat{f}^{[m]}, \theta^{[m]}$$

# Component-Wise Boosting: The Idea



$$\text{Iteration 1: } \hat{f}^{[1]}(x) = \beta b_3(x_3, \theta^{[1]})$$

$$\text{Iteration 2: } \hat{f}^{[2]}(x) = \beta b_3(x_3, \theta^{[1]}) + \beta b_2(x_3, \theta^{[2]})$$

$$\text{Iteration 2: } \hat{f}^{[3]}(x) = \beta b_3(x_3, \theta^{[1]}) + \beta b_2(x_3, \theta^{[2]}) + \beta b_1(x_1, \theta^{[3]})$$

$$\Rightarrow \hat{f}^{[3]}(x) = \beta \left( b_3(x_3, \theta^{[1]} + \theta^{[2]}) + b_1(x_1, \theta^{[3]}) \right)$$

# Component-Wise Boosting: The Algorithm

**Result:** Component-wise boosting model  $\hat{f}(x)$

Initialize  $\hat{f}^{[0]}(x) = \arg \min_{c \in \mathbb{R}} \mathcal{R}_{\text{emp}}(c)$  ;

**for**  $m \in \{1, \dots, M\}$  **do**

    // Update pseudo residuals:

$$r^{[m](i)} = - \left[ \frac{\delta}{\delta f(x^{(i)})} L \left( y^{(i)}, f(x^{(i)}) \right) \right]_{f=f^{[m-1]}}, \quad \forall i \in \{1, \dots, n\} ;$$

    // Get index  $j^*$  of  $m$ -th base-learner from optimizer:

**for**  $j \in \{1, \dots, J\}$  **do**

        // Fit each base-learner  $b_j^{[m]}$  to the pseudo residuals:

$$\hat{\theta}_j^{[m]} = \arg \min_{\theta_j} \sum_{i=1}^n \left( r^{[m](i)} - b_j^{[m]}(x^{(i)}, \theta_j) \right)^2 ;$$

        // Calculate the SSE of the fitted base-learner:

$$\text{SSE}_j = \sum_{i=1}^n \left( r^{[m](i)} - b_j^{[m]}(x^{(i)}, \hat{\theta}_j) \right)^2 ;$$

**end**

    // Add selected component to model:

$$\hat{f}^{[m]}(x) = \hat{f}^{[m-1]}(x) + \beta b_{j^*}^{[m]} \left( x, \theta_{j^*}^{[m]} \right)$$

**end**

**Returns:**  $\hat{f}(x) = \hat{f}^{[m]}(x)$ ;

- Tree-based implementations:
  - `xgboost`
  - `catboost`
  - `gbm`
- Model-based implementations:
  - `mboost` (`gamboost`, `gamboostLSS`)

So, why another boosting implementation?

## **Compboost vs Other Implementations**

---




# Compboost vs Other Implementations

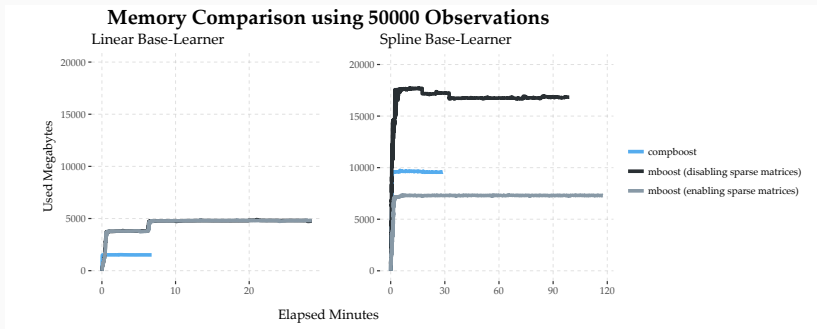
- Compboost is not designed to have huge predictive power such as `xgboost`.
- Hence, the only competitor is `mboost`:
  - `compboost` is not as comprehensive as `mboost` (just two base-learners and three pre defined losses).
  - `compboost` is not able to use sparse matrices at the moment.

# Runtime Comparison

## Benchmark for Increasing Number of Rows



# Memory Comparison



## About Comboost

---

- Installation

```
devtools::install_github("schalkdaniel/compboost")  
library(compboost)
```

# Comboost Members and Member Functions

- **Member Functions:**

- `addBaselearner()`
- `addLogger()`
- `train()`
- `coef()`
- `predict()`
- `risk()`
- `selected()`
- `plot()`
- `...`

- **Public Members:**

- `model`
- `bl.factory.list`
- `loss`
- `optimizer`
- `...`

## Small Usecase

---

# Initializng Model

```
mtcars$mpg_cat = ifelse(mtcars$mpg > 15, "A", "B")
cboost = Comboost$new(mtcars, "mpg", loss = QuadraticLoss$new())

cboost$addBaselearner("wt", "spline", PSplineBlearnerFactory,
  degree = 3, knots = 10, penalty = 2, differences = 2)
cboost$addBaselearner("mpg_cat", "linear", PolynomialBlearnerFactory,
  degree = 1, intercept = FALSE)

cboost$train(2000, trace=FALSE)
cboost

## Componentwise Gradient Boosting
##
## Trained on mtcars with target mpg
## Number of base-learners: 3
## Learning rate: 0.05
## Iterations: 2000
## Offset:20.090625
##
## QuadraticLoss Loss:
##
## Loss function:  $y = (y - f(x))^2$ 
##
##
```




# Plot Results

With `plot()` it is possible to illustrate a specific effect. Additionally, we can specify which iterations we want to visualize. The returned object is an ordinary `ggplot` object:

```
library(ggplot2)
library(ggthemes)

cboost$plot("wt_spline", iters = c(100, 500, 1000, 2000)) +
  labs(title = "Effect of Weight",
       subtitle = "Additive contribution of linear predictor") +
  theme_tufte() +
  scale_color_brewer(palette = "Spectral")
```

# Plot Results



# Using a Custom Logger

1. Define a custom loss which returns the AUC as “loss”:

```
# Define custom "loss function":
aucLoss = function (truth, response) {
  # Convert response on f basis to probs using sigmoid:
  probs = 1 / (1 + exp(-response))

  # Calculate AUC:
  mlr:::measureAUC(probabilities = probs, truth = truth,
    negative = -1, positive = 1)
}

# Define also gradient and constant initialization since they are
# required by the custom constructors:
gradDummy = function (truth, response) { return (NA) }
constInitDummy = function (truth, response) { return (NA) }

# Define loss:
auc.loss = CustomLoss$new(aucLoss, gradDummy, constInitDummy)
```

# Using a Custom Logger

## 2. Register a new out of bag risk logger with the custom loss:

```
cboost$addLogger(logger = OobRiskLogger, use.as.stopper = FALSE,
  logger.id = "auc_oob", auc.loss, 0.01, cboost$prepareData(mtcars[idx.test, ]),
  mtcars[idx.test, "mpg_bin"])
cboost$addLogger(logger = TimeLogger, use.as.stopper = FALSE,
  logger.id = "time", max.time = 0, time.unit = "microseconds")

cboost$train(1000)
```

Iteration	Out of Bag Risk	microseconds
1/10	0.91	1
2/10	0.91	2000
3/10	0.91	3490
4/10	0.91	5218
5/10	0.91	6293
6/10	0.91	7152

## Next Steps

---

- Implementing more base-learner and loss functions.
- Support for multiclass classification.
- Parallel computations.
- Using sparse matrices to store, e.g., spline data matrices.

**Questions?**